

Benchpark and Ramble Tutorial



Ramble

HPCIC Tutorial Series
September 3, 2025



Olga Pearce, LLNL

Doug Jacobsen, Google

Prepared by LLNL under Contract DE-AC52-07NA27344.

Tutorial Presenters



Olga Pearce
LLNL



Doug Jacobsen
Google

Tutorial Agenda (approximate)

Welcome and Introduction	All	10 min
Benchpark Overview and Basics Hands on: Running an experiment, analyzing performance of a scaling study	Olga	80 min
Ramble Overview and Basics + Hands on	Doug	40 min
Benchpark System and Experiment Specifications Hands on: Writing an experiment in Benchpark	Olga	60 min
Q&A and Wrap-up	All	10 min



Tutorial Materials

Find these slides and associated scripts here:

software.llnl.gov/benchpark/basic-tutorial.html

We also have a channel room on Spack slack.

You can join here:

slack.spack.io

Join the **#benchpark-support** channel!

You can continue to ask questions here after the tutorial has ended.

Tutorial Instances: <http://bit.ly/4kGQDlc>



- We have an AWS instance for the hands-on component of this tutorial
- The instance provides:
 - Pre-installed Benchmark and required dependencies
 - Job scheduler

When logging in to the instance:

- Please use a unique username to avoid resource allocation conflicts
 - First initial followed by last name (e.g., John Doe would be jdoe)
- PW: hpctutorial25

We benchmark HPC systems for many reasons

- Procurements
 - Communicate datacenter workload to vendors
 - Co-design systems, monitor progress
 - System acceptance (contractual specification)
- Validation of software stack, tools
 - Compilers
 - Debuggers
 - Correctness tools
 - Performance tools
- Research
 - Programming models
 - Computational methods
 - Performance across architectures



Benchmarking is challenging

- What are we trying to characterize?
- Are we capturing the best the system can do?
- Is something else impacting performance?
- Did we build and run the code in the optimal and reproducible way?

Source: <https://www.top500.org/statistics/perfdevel/>



HPC benchmarks run on diverse HPC hardware



Lawrence Berkeley Nat'l Lab
AMD Zen + NVIDIA



F u g a k u
RIKEN Fujitsu **ARM a64fx**



Lawrence Livermore Nat'l Lab
IBM Power9 + NVIDIA



Lawrence Livermore Nat'l Lab
AMD Zen + Radeon



Oak Ridge National Lab
AMD Zen + Radeon



Argonne National Lab
Intel Xeon + Xe

- Benchmark source code
 - Abstraction (OpenMP, RAJA, Kokkos)
 - Hardware-specific (CUDA, ROCm)
- Optimized code for the CPU and GPU
 - Must make effective use of the hardware
 - Can make 10-100x performance difference
- Rely heavily on system packages
 - Need to use optimized communication and MPI libraries that come with machines

Writing benchmark source code is only the beginning


State-of-the-practice: HPC system benchmarking is manual!

- Building on each system is different, porting the builds to new systems is manual
- Running on each system is different, porting run scripts to new systems is manual
- Systems keep changing, requiring updates to how we build and run benchmarks
- Triggering builds and runs is manual: benchmark results don't stay up to date
- Performance analysis of results is manual



HPC benchmarks are HPC software

- Portability
- Maintenance
- Testing/CI
- Verification
- Reproducibility



All components must work for your system, focus on explainable performance

Benchmarkpark enables complete specification of HPC benchmarks



Infrastructure-as-code benchmark specification codifies:


- Benchmark build and run instructions
- HPC Systems
- HPC Experiments

Leverage advances in HPC automation

- Source code
- Build specification
- Run specification
- CI



github.com/llnl/benchmarkpark



Ramble



Spack

Every part of the specification is codified: used to communicate, automate

Benchmark enables reproducible specifications of benchmarks

Specify

- How to build and run benchmarks on a system

Run

- Run an experiment on a system

Reproduce

- Re-run an experiment on a system

Replicate

- Run an experiment on a new system

Maintain

- CI: Run experiment on HPC systems

Record

- Perf.measurements + *full spec* of experiment

Full specification enables reproducibility, replicability, and automation

HPC System definition for *Performance*

- Resources (CPU cores, GPUs)
- Scheduler (Slurm, flux)
- Process mapping (cores, sockets, GPUs, NICs)
- On-node parallelism (CUDA, ROCM, OpenMP)
- Software stack
 - Compilers (which is best?)
 - MPI implementations (best?)
 - Math libraries



- v
 - systems
 - > all_hardware_descriptions
 - > aws-pcluster
 - > generic-x86
 - > lanl-venado
 - > llnl-cluster
 - > **llnl-elcapitan**
 - > llnl-sierra
 - > riken-fugaku

```

class LlnlElcapitan(System):
    variant(
        "rocm",
        default="5.5.1",
        values=("5.4.3", "5.5.1", "6.2.4"),
        description="ROCm version",
    )
    variant(
        "compiler",
        default="cce",
        values=("gcc", "cce", "rocmcc"),
        description="Which compiler to use",
    )
    def initialize(self):
        super().initialize()
        self.scheduler = "flux"
        self.sys_cores_per_node = "128"
        self.sys_gpus_per_node = "4"

```

benchpark system init --dest=elcap llnl-elcapitan rocm=6.4.2 compiler=cce

BASICS

For the Impatient

Getting Started

Benchpark Commands

Benchpark Workflow

Frequently Asked Questions

CATALOGUE

System Specifications

Benchmarks and Experiments

TUTORIALS

Hello Benchpark Example

Running on an LLNL System

Comparing two Experiments Within Benchpark

USING BENCHPARK

Setting Up a Benchpark Workspace

Building an Experiment in Benchpark

Running an Experiment in Benchpark

Experiment pass/fail and FOMs

Canned Analysis for Scaling Studies

Benchpark Modifiers

System Specifications

The table below provides a directory of information for systems that have been specified in Benchpark. The column headers in the table below are available for use as the `system` parameter in `benchpark setup`.

Search:

name	integrator.vendor	integrator.name	processor.vendor	processor.name	processor.ISA	processor.w
Atos-zen2-A100-Infiniband	Atos	XH2000	AMD	EPYC-Zen2	x86_64	zen2
AWS_PCluster-zen-EFA	AWS	ParallelCluster	AMD	EPYC-Zen	x86_64	zen
DELL-cascadelake-InfiniBand	DELL		Intel	Xeon6248R	x86_64	cascadelake
DELL-sapphirerapids-OmniPath	DELLEMC	PowerEdge	Intel	XeonPlatinum8480	x86_64	sapphirerapi
Fujitsu-A64FX-TofuD	Fujitsu	FX1000	Fujitsu	A64FX	Armv8.2-A-SVE	aarch64
HPECray-haswell-P100-Infiniband	HPECray		Intel	Xeon-E5-2650v3	x86_64	haswell
HPECray-neoverse-H100-Slingshot	HPECray	EX254n	NVIDIA	Grace	Armv9	neoverse
HPECray-zen2-Slingshot	HPECray		AMD	EPYC-7742	x86_64	zen2
HPECray-zen3-MI250X-Slingshot	HPECray	EX235a	AMD	EPYC-Zen3	x86_64	zen3
HPECray-zen4-MI300A-Slingshot	HPECray	EX255a	AMD	EPYC-Zen4	x86_64	zen4
IBM-power9-V100-Infiniband	IBM	AC922	IBM	POWER9	ppc64le	power9
Penguin-icelake-OmniPath	PenguinComputing	RelionCluster	Intel	XeonPlatinum924248C	x86_64	icelake
Supermicro-icelake-OmniPath	Supermicro		Intel	XeonPlatinum8276L	x86_64	icelake
x86_64					x86_64	

HPC systems in Benchmark: July 2025

- 4 in Europe
- 6 in US labs
- 1 in Japan
- 1 at a university
- 2 cloud systems

 all_hardware_descriptions
 aws-pcluster
 common
 csc-lumi
 cscs-daint
 cscs-eiger
 generic-x86
 jsc-juwels
 lanl-venado
 llnl-cluster
 llnl-elcapitan
 llnl-sierra
 riken-fugaku

HPC Experiment definition for *Performance*

- On-node parallelism (CUDA, ROCM, OpenMP)
- Problem sizes
 - Overall problem size, or
 - Per node or per GPU
- Scaling studies
 - How to scale
 - How to decompose
- Resources (cores, GPUs)



Goal: Specify reproducible sets of experiments that map onto specific systems

HPC Experiment in Benchpark: Specify Once

```
class Amg2023(Experiment, OpenMPExperiment, CudaExperiment, ROCmExperiment,
              Scaling(ScalingMode.Strong, ScalingMode.Weak, ScalingMode.Throughput,
                      Caliper)):
    variant(
        "workload",
        default="problem1",
        values=("problem1", "problem2"),
        description="problem1 or problem2",
    )

    def compute_applications_section(self):
        self.register_scaling_config({
            ScalingMode.Strong: {
                "n_resources": lambda var, itr, dim, scaling_factor: var.val(dim) * scale,
                "problem_size": lambda var, itr, dim, scaling_factor: var.val(dim) // scale,
            },
            ScalingMode.Weak: {
                "n_resources": lambda var, itr, dim, scaling_factor: var.val(dim) * scale,
                "problem_size": lambda var, itr, dim, scaling_factor: var.val(dim),
            },
            ScalingMode.Throughput: {
                "n_resources": lambda var, itr, dim, scaling_factor: var.val(dim),
```

```
benchpark experiment init --dest=amg amg2023 +rocm +strong workload=problem2 caliper=mpi,time --system=<system>
```

GETTING STARTED

1. Getting Started with Benchpark
2. Searching Benchpark
3. Editing the experiment (optional)
4. Setting up Benchpark
5. Building the experiment
6. Running an Experiment in Benchpark
7. Analyzing Experiments in Benchpark

FAQ

- What to rerun after edits
- Spack/Ramble versions in Benchpark
- Benchmark not yet in Spack/Ramble

CATALOGUE

System Specifications

Benchmarks and Experiments

Benchpark Help Menu

CONTRIBUTING

Benchmarks and Experiments

	hpl	hpcg	qws
application-domain	['synthetic']	['synthetic']	['qcc
benchmark-scale	['large-scale']	['large-scale']	['we
communication	['mpi']	['mpi']	['mp
communication-performance-characteristics	['network-collectives', 'network-point-to-point']	['network-point-to-point']	[]
compute-performance-characteristics	[]	[]	[]
math-libraries	['blas']	[]	[]
memory-access-characteristics	[]	[]	[]
mesh-representation	[]	[]	[]
method-type	['dense-linear-algebra', 'solver']	['conjugate-gradient', 'solver', 'sparse-linear-algebra']	[]
programming-language	['c']	['c++']	['c+h
programming-model	[]	['openmp']	[]

[← Previous](#)[Next →](#)



Experiments in Benchpark: July 2025

▼ experiments

- > ad
- > amg2023
- > babelstream
- > genesis
- > gpcnet
- > gromacs
- > hpcg
- > hpl
- > ior
- > kripke
- > laghos
- > lammps
- > md-test
- > osu-micro-benchmarks
- > phloem
- > quicksilver
- > qws
- > raja-perf
- > remhos
- > salmon-tddft
- > saxpy
- > smb
- > stream


- HPL, HPCG
- 4 microbenchmarks
- 3 MPI benchmarks
- 8 US, 1 Europe, 2 Japan

`benchpark analyze` for generating pre-defined analysis charts



See <https://software.llnl.gov/benchpark/benchpark-analyze.html>

Contributions from 11 organizations (60% non-LLNL)





Public CDash
dashboard for
CI tests

Nightly [benchpark-develop] 30 builds

[\[view timeline\]](#)

Site	Build Name	Update	Test			Start Time ▾
		Revision	Not Run	Fail	Pass	
dane6	dane/kripke caliper=mpl scaling=weak scaling-iterations=2 compiler=oneapi	74bb04	0	0 ₋₁	1 ⁺¹	4 hours ago
dane1	dane/kripke caliper=mpl scaling=strong scaling-iterations=2 compiler=oneapi	74bb04	0	0	1	4 hours ago
dane16	dane/laghos caliper=mpl scaling=strong scaling-iterations=2	74bb04	0	0	1	4 hours ago
dane1	dane/stream	74bb04	0	0	1	4 hours ago
dane7	dane/smb	74bb04	0	0	1	5 hours ago
dane3	dane/remhos	74bb04	0	0	1	5 hours ago
dane6	dane/phloem	74bb04	0	0 ₋₁	1 ⁺¹	5 hours ago
dane3	dane/osu-micro-benchmarks	74bb04	0	0	1	5 hours ago
dane6	dane/laghos	74bb04	0	0	1	5 hours ago
dane4	dane/kripke	74bb04	0	0	1	5 hours ago

Tuolumne HPECray-zen4-MI300A-Slingshot [benchpark system init --dest=tuolumne-system lnl-elcapitan cluster=tuolumne] 30 builds

[\[view timeline\]](#)

Site	Build Name	Update	Test			Start Time ▾
		Revision	Not Run	Fail ▾	Pass	
tuolumne2150	tuolumne/kripke +rocm scaling=strong scaling-iterations=2 gpumode=CPX	9edf8b	0	1 ⁺¹	0 ₋₁	10 hours ago
tuolumne2149	tuolumne/raja-perf +rocm	e1a1c5	0	0	1	12 minutes ago
tuolumne2149	tuolumne/kripke +rocm scaling=strong scaling-iterations=2 gpumode=TPX	e1a1c5	0	0	1	19 minutes ago
tuolumne2149	tuolumne/amg2023 +rocm	e1a1c5	0	0	1	20 minutes ago
tuolumne2149	tuolumne/laghos +rocm scaling=strong scaling-iterations=2 gpumode=TPX	e1a1c5	0	0	1	26 minutes ago
tuolumne2150	tuolumne/laghos +rocm scaling=strong scaling-iterations=2 gpumode=CPX	e1a1c5	0	0	1	27 minutes ago
tuolumne2149	tuolumne/laghos +rocm	e1a1c5	0	0	1	28 minutes ago
tuolumne1001	tuolumne/amg2023 +rocm scaling=strong scaling-iterations=2 gpumode=TPX	e1a1c5	0	0	1	38 minutes ago
tuolumne2149	tuolumne/kripke +rocm scaling=strong scaling-iterations=2 gpumode=CPX	e1a1c5	0	0	1	39 minutes ago
tuolumne2149	tuolumne/kripke +rocm	e1a1c5	0	0	1	40 minutes ago

Benchmark codifies benchmarking steps

- Benchmark does **not** replace benchmark source, build system, or Spack package
- Benchmark manages benchmark **experiments** and how they map onto **systems** with specified (or default)
 - Compilers
 - MPI/comm libraries
 - Math libraries
- Start running benchmark experiments on your system with just a few commands

 `git clone git@github.com:LLNL/benchmark.git`

 `benchmark list systems`
`benchmark system init --dest=elcap llnl-elcapitan`

`benchmark list experiments`
 `benchmark experiment init --system=elcap --dest=amg`
`amg2023 +rocm +strong`

 `benchmark setup amg workspace`
`ramble workspace setup`

 `ramble on`

Who does Benchmark target

People who want to use or distribute benchmarks for HPC!

1. End Users of HPC Benchmarks

- Install, run, analyze performance of HPC benchmarks

2. Benchmark Developers

- People who want to share their benchmarks

3. Procurement teams at HPC Centers


- Curate workload representation, evaluate and monitor system progress

4. HPC Vendors

- Understand the curated workload of HPC centers, propose systems

Catalogued library of *working* benchmarks

- Enables exploration of large configuration space
 - Architectural
 - Software stack
 - Temporal
- What architecture and system configuration is best for my benchmark?
- On my system, is OpenMPI good enough?
- What purpose will Benchpark help you address?




Building a community to contribute/benefit

Benchmark: Open collaborative repository for reproducible specifications of HPC benchmarks

Infrastructure-as-code benchmark specification

enables **reproducibility, replicability, and automation**

- HPC Systems
- HPC Experiments



- Tagging, keywords for publications
- Performance metrics, metrics of usefulness
- Dashboards: Archive specs+results

- CI pipelines on PRs from GitHub at data centers across the world and in the cloud

*Olga Pearce et al, Continuous Benchmarking, HPCTests SC|23

*Olga Pearce et al, Repeat, Replicate, Reproduce, ACM REP 2025

Full specification enables reproducibility, replicability, and automation

Benchpark roadmap and community engagement

Future directions:

- Suite curation: Reproducible specification of an entire suite
- Tagging: Keywords for publications, finding benchmarks
- Metrics: Performance, usefulness
- Dashboards: Archive+share *specs*+results, Slices in configuration space
- CI pipelines at data centers across the world and in the cloud

Community Engagement:

- Co-design / vendors on board, but incentive for app teams? (carrot or stick?)
- Who owns which parts of the specification and approves changes?
- Who finances R&D and maintenance?
- ROI for the people working on it? → think about post docs, researchers, etc.

Collaboration, reproducibility, fully-specified public results

Hands On Session 1

Running an existing benchmark on an existing system

Follow script at software.llnl.gov/benchpark/basic-tutorial.html

Tutorial Instances: <http://bit.ly/4kGQDlc>

- We have an AWS instance for the hands-on component of this tutorial
- The instance provides:
 - Pre-installed Benchpark and required dependencies
 - Job scheduler



When logging in to the instance:

- Please use a unique username to avoid resource allocation conflicts
 - First initial followed by last name (e.g., John Doe would be jdoe)
- PW: hpctutorial25

Tutorial Materials

Find these slides and associated scripts here:

software.llnl.gov/benchpark/basic-tutorial.html

We also have a channel room on Spack slack.

You can join here:

slack.spack.io

Join the **#benchpark-support** channel!

You can continue to ask questions here after the tutorial has ended.

What is Ramble?

What is Ramble?

[Ramble](#) is an open-source experimentation framework written in python.

Ramble's primary goals are to:

- **Accelerate productivity**
- **Improve experiment portability**
- **Increase reproducibility information**
- **Encoding domain knowledge**

Ramble is particularly good at generating parameterized “experiments”, but can be used for more than performance focused experiments.

What is Ramble?


Ramble is:

- Written in python
- Multi-platform
- Heavily based on Spack (Written on top of Spack's infrastructure)
- An experimentation framework (Related to running experiments like Spack is to building software)
- Open Source: <https://github.com/GoogleCloudPlatform/ramble>

Ramble provides:

- A flexible templating engine
- A domain specific language for describing an application and its workloads
- Standardized definitions of how workloads are configured and executed

Ramble Architecture



Ramble Components

Object Definitions

Portable definitions written in python, that represent:

Applications:

- Executables (commands to run)
- Input files (things to download)
- Software (things to install)
- Workloads (executables + input files)
- Figures of merit (things to analyze)

Package Managers (spack, eessi)

Workflow Managers (slurm, batch, gke)

Modifiers (lscpu, intel-aps, etc..)

Ramble supports **public** and **private** object definitions through configurable repos.

Does not include experiment or system details.

Experiment Configuration (Workspace)

Self contained directory describing a set of **parameterized** experiments to execute.

Contains:

- YAML File, describes:
 - Experiments
 - Software stack
 - Execution workflow
- Template files

After setup, contains:

- Input files
- Experiment execution directories
- Software environments

Experiment Directory

Each experiment has **their own** execution directory.

These contain:

- Rendered versions of the template files (which can then be executed)

After executing the might contain:

- Links to the input files
- Output files from the experiment to be analyzed

Ramble Components

Software Stack

Ramble supports using a predefined software stack, or generating a new software stack to produce experiments.

Supported software stack workflows in Ramble include:

- Pre-installed software (Considered “Existing”)
 - e.g. ISV applications installed in a static location
 - Software installed manually through another mechanism (yum, apt, pip, etc..)
- Driving the creation of a software stack:
 - Using a 3P package manager (currently supports [spack](#), [pip](#), [eessi](#), [environment-modules](#))
 - Can parameterize various aspects of the stack
 - Can use a binary cache to accelerate deployments, and provide binary objects to customers

When Ramble manages the software stack, it can parameterize portions of the stack to generate various studies. This includes:

- Changing compilers
- Changing dependencies (e.g. MPI impl., CUDA vs. ROCM)
- Package versions (e.g. CUDA versions)
- Optimization Targets (e.g. Icelake vs. Zen 3)
- etc...

Allows very flexible descriptions of the software environment

Spack is used to distribute [well defined software stacks](#).

Various organizations provide public spack caches, which can make it faster to deploy a complex software stack.


How does Ramble work?

What is Ramble's Workflow?

Ramble Controlled Step

User Input Step

Output Step



Ramble Workspace:

A workspace is a self contained directory, that contains:

- software environments
- input files
- experiments

Workspaces are independent, and should represent a set of independent experiments you want to execute

Workspace: my_experiments

```
configs:  
└─ ramble.yaml  
   execute_experiment.tpl
```

inputs

software

experiments

Workspace Config:

Ramble's workspace config YAML can generate many experiments with only a little syntax.

applications:

wrfv3: Application Definition Name

n_repeats: 5 } Define 5 repeats per experiment

variables:

processes_per_node: [30, 56]
partition: ['c2', 'c2d'] } Define two vectors, of length 2.

n_ranks: '{processes_per_node}*{n_nodes}'

workloads:

CONUS_2p5km: Workload Name (from Application Definition)

experiments:

scaling_study_{partition}_{nodes}nodes:

variables:

n_nodes: [1, 2, 4, 8, 16, 32] } Define one vector, of length 6

matrix:
- n_nodes } Define one matrix, using the n_nodes vector. Shape is 1x6.

- Vectors that are not used by a matrix are zipped together (must be the same length)
- Zip of vectors are crossed with any matrices
- Result is: 1x6 (matrix) x 2 (vectors) x 5 (repeats) = 60, where each index is a 4-element tuple: (processes_per_node, n_ranks, n_nodes, repeat_idx)

Workspace Config:

In addition to the YAML config, Ramble has a template engine to render scripts:

```
#!/bin/bash
```

```
{command}
```



```
#!/bin/bash
```

```
source <path/to/spack>  
spack env activate <path/to/env>  
cp <inputs>/* <experiment_dir>/.  
mpirun -n <n_ranks> wrf.exe
```

Workspace Setup:

Workspace: wrf-demo

```
configs:  
  - ramble.yaml  
  - execute_experiment.tpl
```

inputs

software

experiments

Execute input phases:

- Determine which inputs are necessary
- Download them
- Define variables for referring to these inputs

Execute software phases (Only used on SpackApplications currently):

- Install necessary compilers
- Create spack environments for the experiments
- Using spack, install the required software
- Define variables for referring to the spack environments

Execute experiment creation phases:

- Create experiment execution directories
- Define variables specific to the experiment
- Render any *.tpl files into the execution directories
- Append the experiment to \$workspace/all_experiments

Execute Experiments:

Workspace: wrf-demo

```
configs:  
└── ramble.yaml  
    execute_experiment.tpl
```

```
software/  
└── wrfv3 .CONUS_2p5km  
    spack.yaml
```

```
inputs:  
└── ...
```

all_experiments

```
experiments/  
└── wrfv3  
    └── CONUS_2p5km  
        ├── scaling_study_c2_16nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2_1nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2_2nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2_32nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2_4nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2_8nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2d_16nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2d_1nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2d_2nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2d_32nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        ├── scaling_study_c2d_4nodes  
        │   ├── execute_experiment  
        │   └── execute_experiment  
        └── scaling_study_c2d_8nodes  
            ├── execute_experiment  
            └── execute_experiment
```

After setting up a workspace, the experiments can be executed using:

- `$workspace/all_experiments`
- `ramble on` (with an activated workspace)

Depending on the `ramble.yaml` and `execute_experiment.tpl` this:

- Executes experiments sequentially
- Submits experiments to a workload manager

Analyze Workspace:

After experiments are executed, Ramble can extract their figures of merit.

This is done through:

```
ramble workspace analyze
```

Ramble will process the output files described in an experiment's application.py, and extract success criteria and figures of merit.

These are then written to a file (e.g. results.txt) in the provided format (can be text, yaml, or json).



What other things can you do with Ramble?

Parameterizing Software Definitions

A common workflow in optimizing software is exploring how modifications to a software stack impact performance. Ramble's support of these workflows can be seen below:

```
ramble:
```

```
...
applications:
  wrfv3:
    workloads:
      CONUS_2p5km:
        experiments:
          '{mpi}_{partition}_{n_nodes}nodes':
            variables:
              n_nodes: [1, 2, 4, 8, 16, 32]
              mpi: ['impi', 'ompi']
              env_name: 'wrf-{mpi}'
            matrix:
              - n_nodes
              - mpi
```

```
[ramble:]
```

```
...
software:
  packages:
    wrf:
      pkg_spec: wrf@3.9.1.1
    impi:
      pkg_spec: intel-oneapi-mpi@2021.9
    ompi:
      pkg_spec: openmpi@4.1.4
  environments:
    wrf-{mpi}:
      packages:
        - wrf
        - '{mpi}'
```

Parameterizing Software Definitions

A common workflow in optimizing software is exploring how modifications to a software stack impact performance. Ramble's support of these workflows can be seen below:

ramble:

```
...
applications:
  wrfv3:
    workloads:
      CONUS_2p5km:
        experiments:
          '{mpi}_{partition}_{n_nodes}nodes':
            variables:
              n_nodes: [1, 2, 4, 8, 16, 32]
              mpi: ['impi', 'ompi']
              env_name: 'wrf-{mpi}'
            matrix:
              - n_nodes
              - mpi
```

[ramble:]

```
...
software:
  packages:
    wrf:
      pkg_spec: wrf@3.9.1.1
      impi:
      pkg_spec: intel-oneapi-mpi@2021.9
      ompi:
      pkg_spec: openmpi@4.1.4
    environments:
      wrf-{mpi}:
        packages:
          - wrf
          - '{mpi}'
```

Defines spack environments

Variable
expansion

MPI
packages

Performance Analysis Tools

Modifiers all composable definitions to edit the experiment behavior

ramble:

```
...
modifiers:
- name: intel-aps
applications:
  wrfv3:
    workloads:
      CONUS_2p5km:
        experiments:
          '{mpi}_{partition}_{n_nodes}nodes':
            variables:
              n_nodes: [1, 2, 4, 8, 16, 32]
              mpi: ['impi', 'ompi']
              env_name: 'wrf-{mpi}'
            matrix:
              - n_nodes
              - mpi
```

[ramble:]

```
...
software:
  packages:
    wrf:
      pkg_spec: wrf@3.9.1.1
    impi:
      pkg_spec: intel-oneapi-mpi@2021.9
    ompi:
      pkg_spec: openmpi@4.1.4
  environments:
    wrf-{mpi}:
      packages:
        - wrf
        - '{mpi}'
```


Performance Analysis Tools

Modifiers all composable definitions to edit the experiment behavior

ramble:

```
...
modifiers:
- name: intel-aps ←
applications:
  wrfv3:
    workloads:
      CONUS_2p5km:
        experiments:
          '{mpi}_{partition}_{n_nodes}nodes':
            variables:
              n_nodes: [1, 2, 4, 8, 16, 32]
              mpi: ['impi', 'ompi']
              env_name: 'wrf-{mpi}'
            matrix:
              - n_nodes
              - mpi
```

[ramble:]

```
...
software:
  packages:
    wrf:
      pkg_spec: wrf@3.9.1.1
    impi:
      pkg_spec: intel-oneapi-mpi@2021.9
    ompi:
      pkg_spec: openmpi@4.1.4
  environments:
    wrf-{mpi}:
      packages:
        - wrf
        - '{mpi}'
```

Performance Analysis
Tool

Three Rs in HPC Benchmarking

1. Repeat
2. Replicate
3. Reproduce

Three Rs in HPC Benchmarking

1. Repeat: run on same system

2. Replicate

3. Reproduce

Component	Repeat
(1) System H/W	✓
(2) System S/W	✓
(3) Application	✓

- Verifying benchmark still builds and runs
- Handing over benchmarking tasks from one user to another
- Evaluating benchmark performance changes as system ages
- Running multiple trials of a benchmark for noise variability

Three Rs in HPC Benchmarking

1. Repeat

2. Replicate: run with a different software stack

3. Reproduce

Component	Repeat	Replicate
(1) System H/W	✓	✓
(2) System S/W	✓	≈
(3) Application	✓	✓

- Verifying benchmark still builds and runs with update system software
- Validating performance of updated system software
- Evaluating performance differences of different software substitutes (e.g., gcc or icc compilers)

Three Rs in HPC Benchmarking


1. Repeat
2. Replicate
3. Reproduce: run on a different system

Component	Repeat	Replicate	Reproduce
(1) System H/W	✓	✓	X
(2) System S/W	✓	≈	X
(3) Application	✓	✓	≈

- Verifying benchmark builds and runs on diverse hardware
- Evaluating benchmark performance across different hardware
- Exploring potential benefits of hardware upgrades

Benchmark Specification for Three Rs

	Component	Repeat	Replicate	Reproduce
(1)	System H/W	✓	✓	X
(2)	System S/W	✓	≈	X
(3)	Application	✓	✓	≈



Benchmark Specification for the Three Rs

	Component	Repeat	Replicate	Reproduce
(1)	System H/W	✓	✓	X
(2)	System S/W	✓	≈	X
(3)	Application	✓	✓	≈

Repeat as
a different
user

```
1 sysa:~ user1$ benchpark system init --dest=sysa intel-cluster
  cluster=sysa
2 sysa:~ user1$ benchpark experiment init --dest=qs
  quicksilver +openmp +weak
3 sysa:~ user1$ benchpark setup ./qs ./sysa workspace/
```

```
1 sysa:~ user2$ benchpark system init --dest=sysa intel-
  cluster cluster=sysa
2 sysa:~ user2$ benchpark experiment init --dest=qs
  quicksilver +openmp +weak
3 sysa:~ user2$ benchpark setup ./qs ./sysa workspace2/
```

Replicate
with
different
SW stack

```
1 sysa:~ user1$ benchpark system init --dest=sysa1 intel-
  cluster cluster=sysa compiler=gcc1
2 sysa:~ user1$ benchpark experiment init --dest=qs
  quicksilver +openmp +weak
3 sysa:~ user1$ benchpark setup ./qs ./sysa1 workspace/
```

```
1 sysa:~ user2$ benchpark system init --dest=sysa2 intel-
  cluster cluster=sysa compiler=gcc2
2 sysa:~ user2$ benchpark experiment init --dest=qs
  quicksilver +openmp +weak
3 sysa:~ user2$ benchpark setup ./qs ./sysa2 workspace2/
```


Reproduce
on a
different
system

```
1 sysa:~ user1$ benchpark system init --dest=sysa intel-
  cluster cluster=sysa
2 sysa:~ user1$ benchpark experiment init --dest=qs-openmp
  quicksilver +openmp +weak
3 sysa:~ user1$ benchpark setup ./qs-openmp ./sysa worksp/
```

```
1 sysb:~ user2$ benchpark system init --dest=sysb ibm-cluster
2 sysb:~ user2$ benchpark experiment init --dest=qs-cuda
  quicksilver +cuda +weak
3 sysb:~ user2$ benchpark setup ./qs-cuda ./sysb work2/
```

Repeat runs of Quicksilver on Ruby using *same compiler*


Component	Repeat	Replicate	Reproduce
(1) System H/W	✓	✓	X
(2) System S/W	✓	≈	X
(3) Application	✓	✓	≈



Benchmark	Quicksilver
Variants	+mpi, +openmp
Num Ranks	4, 8, 12, 16, 20, 24, 28, 32
Num Nodes	1
Scaling	weak
Repetitions	4 (Increase as appropriate for analysis)

Replicate runs of Quicksilver on Ruby with *different compilers*




	Component	Repeat	Replicate	Reproduce
(1)	System H/W	✓	✓	X
(2)	System S/W	✓	≈	X
(3)	Application	✓	✓	≈



Compiler	Version	How Installed
gcc	12.1.1	system-provided
gcc	12.1.0	self-installed

Reproduce runs of Quicksilver across *different systems*

Component	Repeat	Replicate	Reproduce
(1) System H/W	✓	✓	X
(2) System S/W	✓	≈	X
(3) Application	✓	✓	≈




System	P-value	Statistically Significant?
Ruby	0.490	No
Sierra	0.349	No
Dane	0.00000105	Yes


System Name	CPU Hardware	Number of Cores
Ruby	Intel Cascade Lake	56
Sierra	IBM Power 9	44
Dane	Intel Sapphire Rapids	112

*Gap on Dane is statistically significant


Analysis of calltrees to understand performance difference



Self-Installed Compiler



System-Provided Compiler



Takeaways

- Formalize how to repeat, replicate, and reproduce HPC benchmarking experiments
- Showcase how Benchpark enables reproducibility through specification of experiments and HPC systems
- Full record of the specification is preserved for introspection and reproducibility

Hands On Session 2

Adding an experiment for a benchmark

Follow script at <http://software.llnl.gov/benchpark/add-an-experiment.html>

Tutorial Instances: <http://bit.ly/4kGQDlc>

- We have an AWS instance for the hands-on component of this tutorial
- The instance provides:
 - Pre-installed Benchpark and required dependencies
 - Job scheduler



When logging in to the instance:

- Please use a unique username to avoid resource allocation conflicts
 - First initial followed by last name (e.g., John Doe would be jdoe)
- PW: hpctutorial25

Performance Tools: Caliper



- Caliper is an performance profiling library
- Integrates a performance profiler into your program
 - Profiling is always available
 - Simplifies performance profiling for application end users
- Common instrumentation interface
 - Provides program context information for other tools
- Designed for HPC
 - Supports MPI, OpenMP, CUDA, HIP, Kokkos, RAJA
- software.llnl.gov/Caliper

Performance Tools: Thicket



- Thicket is a toolkit for Exploratory Data Analysis
- Enables exploratory data analysis of multi-run data
- Compose data from diff. sources and types
 - Different scaling (e.g., strong, weak)
 - Different application parameters
 - Different compilers and optimization levels
 - Different hardware types (e.g., CPUs, GPUs)
 - Different performance tools
- Perform analysis on the thicket of runs
 - Manipulate the set of data
 - Visualize the dataset
 - Perform analysis on the data
 - Model data
 - Leverage third-party tools in the Python ecosystem

thicket.readthedocs.io

Ramble modifiers in Benchpark encapsulate reusable patterns to perform a specific configuration of an experiment

- Affinity modifier for pinning threads
- Allocation modifier to request resources
- Hwloc modifier to capture hardware topology
- Caliper modifier to profile an application

```
benchpark experiment init --dest=amg amg2023 --system=elcap --rocm workload=problem2 +strong caliper=mpl.time
```




Join us after the tutorial!

Tutorial material: software.llnl.gov/benchpark/basic-tutorial.html

Connect with us on Spack slack

#benchpark-support

We want your
feedback!



Contribute systems, benchmarks, experiments,
and features on GitHub



slack.spack.io
#benchpark-support



★ Star us on GitHub!
github.com/llnl/benchpark





Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.